

# Personalized human video pose estimation

MATLAB code for propagating human pose annotation throughout a video, as detailed in the paper:

J. Charles, T. Pfister, D. Magee, D. Hogg and A. Zisserman "[Personalized human video pose estimation](#)", CVPR 2016.

The code has been tested to work in both Windows 7 and Linux and is also equipped to run across a CPU cluster.

This work is licensed under the Creative Commons Attribution 4.0 International License. To view a copy of this license, visit <http://creativecommons.org/licenses/by/4.0/> or send a letter to Creative Commons, PO Box 1866, Mountain View, CA 94042, USA.

Code is provided by James Charles ([jjcvision@gmail.com](mailto:jjcvision@gmail.com)).

## What it does

Provided with:

- A video of a person,
- Dense optical flow (e.g. by using [DeepFlow](#)),
- A few initial pose annotations obtained either manually or automatically,

this code propagates initial pose annotations across the whole video to be later used for personalizing a generic pose estimator, such as [2].

## What is included

- All stages of propagation, as detailed in [1] are included i.e. spatial matching, temporal propagation and self-evaluation.
- A script for setting up the training material to fine-tune the model of [2] on propagated pose annotations.
- Data for running the demo (360MB): [demo\\_data.zip](#)

## Prerequisites

Install the following prior to running:

- MATLAB 2012a or later (may work on earlier versions but untested)
- MATLAB Image Processing Toolbox
- [VLFeat](#) 0.9.16 or later.
- [Deepflow](#) (used for running on your own video, but not required for the demo).

## Self-contained packages

The following packages are also required but are included with this code so as to be self-contained:

- FLANN <http://www.cs.ubc.ca/research/flann/>
- liblinear <https://www.csie.ntu.edu.tw/~cjlin/liblinear/>

FLANN mex files are pre-compiled. If there are problems using these then please download, install and compile FLANN from: <http://www.cs.ubc.ca/research/flann/>. If you get the MATLAB error: `cannot open with static TLS` when running the demo, restarting MATLAB normally fixes the problem.

## Setup

Compile mex files from within MATLAB:

```
>> compile
```

Setup MATLAB paths:

```
>> set_paths
```

## Running the demo

First download the demo video, optical flow file and initial pose annotations (360MB):  
[demo\\_data.zip](#)

Extract the contents of the zip and retain the folder structure. Initial pose annotations were recovered using the automatic methods described in [1].

### Set paths

Folder options in `./options/part_detector_options.m` need to be set according to your file system before running the demo:

```
folder.main = 'folder where you extracted the demo data';
```

Also, set the experiment folder to somewhere you would like to cache the output. You will require around 600MB. The default is in the current working director.

```
folder.experiment = 'folder where I want to store the cache';
```

### Chose input video

There is the option of running on one of two demo videos. A very short video (10s) useful for testing

the installation: `E7ULR-yfNnk_vshort`. And a longer video (1min) `E7ULR-yfNnk_cut` where more benefit of personalization is observable. In general, the longer the video, the greater the benefit from pose propagation. Set the chosen video in `demo.m`

```
videoname = 'E7ULR-yfNnk_vshort';
```

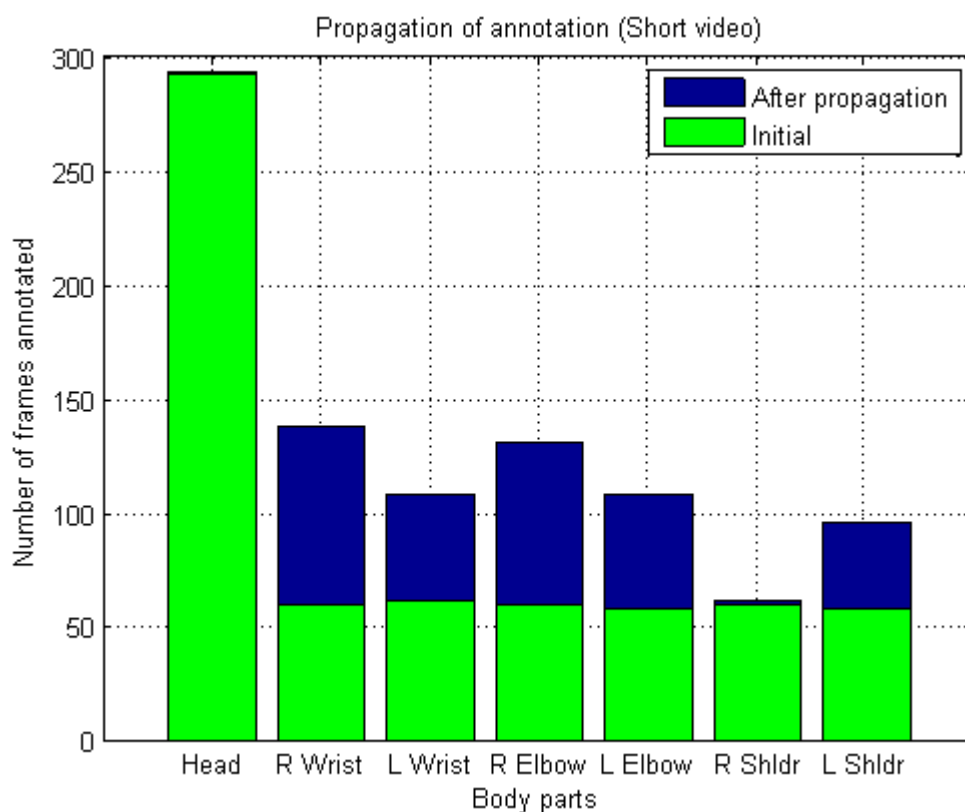
For best pose improvements on the longer video, we recommend setting the optical flow step size `opts.flow.stepsize` in `./options/propagation_options_youtube.m` to 30. When using the shorter video you may leave this set on 3 for speed.


## Running the demo

Run the demo:

```
>> demo
```

First, a visualization of the initial pose estimates on the demo video will play. Afterwards, pose propagation will run for one iteration and then a visualization of propagated poses will be displayed. Depending on the demo video used, one of the following propagation coverage graphs will be displayed:





## Using multiple CPUs

Using one cpu, the short demo will complete in approx. 1 hour, the longer video taking approx. 4-6 hours. On a 12 core machine the longer demo video will complete in approx. 30 minutes. Multiple CPUs can be requested by changing the variable `num_cpus` in `demo.m` provided you have the `Parallel Processing Toolbox` installed. Memory requirements for running the longer demo is around 1.5GB and will increase with the number of cpus requested if running on one machine.

## Running on your own video

Resize the video so that the person has shoulder width approx. 100px wide.

Copy your video into the `videos` folder within the demo, e.g. if your video is called `myvideo.avi`, place this in the folder `./demo_data/videos/`.

Compute dense optical flow across the whole video (we recommend using DeepFlow from <http://lear.inrialpes.fr/src/deepflow/>). Dense optical flow needs to be stored in a `.mat` containing the array `flow` and `minmax`.

A function called `collect_optic_flow` for producing a flow file in this format is provided but requires Deepflow to be installed. Set the path to the compiled binary of Deepflow in `propagation_options_youtube.m`

```
opts.deepflowstatic = '/DeepFlow_release2/deepflow2-static';
```

Compute the required dense optical flow file

```
>> collect_optic_flow('myvideo','youtube');
```

The array `flow` is a `uint8` array of size `frame_height` x `frame_width` x 2 x `N`, where `N` is the number of video frames and the 3rd dimension holds the `u` and `v` flow vector fields, respectively.

The flow vectors are scaled (per frame) between 0 and 255 according to the minimum and maximum values within a frame.

The array `minmax` is a `double` array of size 2 x 2 x `N`. Rows represent `u` and `v` respectively and the columns are the min and max flow values respectively, per frame.

These arrays are saved in `./demo_data/optical_flow/` to a file called `myvideo.mat`

Next, produce initial body joint annotations (2D locations of 7 upper body joints) and save to a MATLAB structure called `detections.manual`:

- `detections.manual.locs` is a 2 x 7 x `N` array representing pose annotations for a selection of `N` frames. A pose for each frame is given by a set of `x` and `y` (row 1 and row 2) body joint locations for the head, right wrist, left wrist, right elbow, left elbow, right shoulder and left shoulder, respectively, per column (person centric).
- `detections.manual.frameids` = 1 x `N` array indicating the annotated frame IDs.

Save this struct:

```
>> save('./demo_data/initialisation/myvideo.mat','detections');
```

If a body joint is unknown for a particular frame, then this can be indicated by setting the `x` and `y` coordinates in `detections.manual.locs` to -999

Finally, initial poses can be propagate for one iteration:

```
>> detections = propagate('myvideo','youtube',...  
    './demo_data/initialisation/myvideo.mat',1,1,1);
```

Pose output is visualized with:

```
>> show_skeleton('./demo_data/videos/myvideo.avi',...  
    1,detections.manual.frameids,detections.manual.locs,0);
```

## Settings

A number of useful parameter settings can be set for adjusting speed of execution and performance. These can be found in the option files `./options/part_detector_options_youtube.m`

and `./options/propagation_options_youtube.m`.

We recommend experimenting with the following:

- `opts.model.forest.numtrees` for setting the number of trees in the random forest part detector. Lower numbers will reduce training time, ensure this number is a multiple of 2.
- `opts.model.rotations` for setting rotation augmentation during part detector training, setting less rotations will reduce memory requirements during training.
- `opts.flow.numneighbours` step size used for selecting frames to temporally propagate (i.e. set to 1 to select every frame or 2 to select every 2nd frame, etc.). Higher values will increase speed.
- `opts.flow.stepsize` temporal window used to propagate forward and backwards in time. In [1], this was set to 30.

## Personalizing the ConvNet of [2]

Personalization is achieved by fine-tuning the ConvNet of [2] (Caffe model) on the propagated pose estimates.

The function `fusion.setupFinetuningCropped` is designed to simplify the task of setting up the training data so the model can be fine-tuned in Caffe.

Various options should be set correctly prior to running this function and a detailed guide is provided:

### Fine-tuning Caffe model of [2]

Download the model and source code of [2] from: <https://github.com/tpfister/caffe-heatmap> and install Caffe.

Set the fine-tuning options in `./options/propagation_options_youtube.m` and provide absolute paths.

Set filename to generic ConvNet model

```
opts.cnn.finetune.model_filename = 'caffe-heatmap-flic.caffemodel';
```

Set folder to where training data will be generated

```
opts.cnn.finetune.ramdisk_folder = '/mnt/ramdisk/data/';
```

Set folder to where snapshots and training scripts are created

```
opts.cnn.finetune.main_save_folder = '/fusion_training/';
```

Set the root folder of the Caffe install

```
opts.cnn.finetune.cafferoot = '/caffe-heatmap/';
```

The `demo.m` script shows an example use of the function `fusion.setupFinetuningCropped`. In `demo.m` set `isFineTune` to `true` and re-run the script. You will then be signaled to run the generated `train.sh` bash script.

Model snapshots are saved in the snapshot folder within `opts.cnn.finetune.main_save_folder`.

**Note.** Body joint locations which are missing from frames are indicated by having the value `-999`. This occurs in frames where all body joints have not been successfully propagated, resulting in parts of poses. Therefore, the Euclidean loss function for [2] should be altered so as to provide zero loss for these missing body joint locations.

## Citation

If you use this code then please cite:

```
@InProceedings{Charles16,  
  author      = "Charles, J. and Pfister, T. and Magee, D. and Hogg, D. and  
Zisserman, A.",  
  title       = "Personalizing Human Video Pose Estimation",  
  booktitle   = "IEEE Conference on Computer Vision and Pattern Recognition",  
  year        = "2016",  
}
```

## References

[1] J. Charles, T. Pfister, D. Magee, D. Hogg and A. Zisserman "[Personalized human video pose estimation](#)", CVPR 2016.

[2] T. Pfister J. Charles and A. Zisserman "[Flowing ConvNets for Human Pose Estimation in Videos](#)", ICCV 2015.